


Programming of Supercomputers

Assignment 1:

Single Node Performance

Prof. Michael Gerndt
Isaias A. Compres Urena

SuperMUC



Introduction

Supercomputing:

- Performance as one of the main goals
- Large amount of parallelism in hardware
 - Multiple sockets
 - Multiple cores
 - SIMD
 - Accelerators
- Large number of nodes aggregated
 - Distributed memory
- Multiple socket nodes and high core counts with SIMD
- ***Specialized networks***
 - Low latency
 - High bandwidth
 - Complex topologies

Course Format

Assignments:

- 4 assignments
- 1 Final summary video

Grading (groups of 2 or 3 students):

- Each assignment and final video will be given a 0-100 grade
 - Linear conversion to German grading
 - 40 : 4,0
 - >93 : 1,0
- These grades will be weighted as follows:
 - Assignment 1: 15%
 - Assignment 2: 15%
 - Assignment 3: 25%
 - Assignment 4: 30%
 - Video: : 15%

Login to SuperMUC, Documentation

- First change the standard password
 - <https://idportal.lrz.de/r/entry.pl>
- Login via
 - lxhalle due to restriction on connecting machines
 - `ssh <userid>@supermuc.lrz.de`
 - No outgoing connections allowed
- Documentation
 - <http://www.lrz.de/services/compute/supermuc/>
 - <http://www.lrz.de/services/compute/supermuc/loadleveler/>
 - Intel compiler:
<http://software.intel.com/sites/products/documentation/hpc/composerxe/en-us/2011Update/cpp/lin/index.htm>

Building the Benchmark

- Load the required modules:

```
module unload mpi.ibm  
module load mpi.intel
```
- Update your Makefile (refer to the provided instructions)
 - Baseline
 - OpenMP
 - MPI
 - MPI + OpenMP
- Build and verify that the binaries were created
- Run the benchmark in the login node
- Identify your **performance metric** from the output!
 - More is better or less is better?
 - Check the benchmark's documentation online

Batch Scripts

- Advantages
 - ***Reproducible performance***
 - Run larger and longer running jobs
- Several job classes available
 - ***Test (recommended for this assignment's tasks)***
 - Phase 1:
 - Max 1 island, 32 nodes, 30 minutes, 1 job in queue
 - Phase 2:
 - Max 1 island, 20 nodes, 30 minutes, 1 job in queue
 - Micro
 - Phase 1:
 - Max 1 island, 32 nodes, 48 hours, 8 jobs in queue
 - Phase 2:
 - Max 1 island, 20 nodes, 48 hours, 8 jobs in queue

Submitting a Batch Job

- `llsubmit ll.sh`
 - Submission to batch system
- `llq -u $USER`
 - Check status of own jobs
- `llcancel <jobid>`
 - Kill job if no longer needed
 - Obtain the <jobid> from the `llq` output

```
#!/bin/bash
#@ wall_clock_limit = 00:20:00
#@ job_name = pos-lulesh-openmp
#@ job_type = MPICH
#@ class = micro
#@ output =
pos_lulesh_openmp_$(jobid).out
#@ error =
pos_lulesh_openmp_$(jobid).out
#@ node = 1
#@ total_tasks = 16
#@ node_usage = not_shared
#@ energy_policy_tag = lulesh
#@ minimize_time_to_solution = yes
#@ island_count = 1
#@ queue

. /etc/profile
. /etc/profile.d/modules.sh

export OMP_NUM_THREADS=16
./lulesh2.
```


Use CPU hours responsibly

- Specify job execution as tight as possible
 - In this assignment, 10 minutes is sufficient
- Only request the number of nodes required.
 - 1 node is sufficient for all tasks in assignment 1.
- Small tests can be done in the login node
 - Create a batch only after you are ready to collect results
 - Running in a batch eliminates interference from other users.
- All types of runs can be tested in the login node
 - Baseline
 - MPI,
 - OpenMP and
 - Hybrid

Assignment 1: Single Node Performance

Single-thread Performance

- GPROF
 - Flat profile
 - Call graph
- Compiler Flags
 - GCC
 - ICC
- Optimization Pragmas
 - GCC pragmas
 - ICC pragmas

Multi-thread Performance

- OpenMP
 - Single process
 - Scaling with threads
 - Shared address space
 - Direct load and stores
 - Coherency, locks, etc.
- MPI
 - Multiple processes
 - Only certain process counts valid
 - Scaling with processes
 - Separate address spaces
 - Messages
- MPI + OpenMP